

Supervised Regression (ANNs)

Presented by: Reem Elmahdi

Agenda

- ❑ Supervised Regression (SR) definition
- ❑ ML models used in SR problems
- ❑ Biological Neural Networks
- ❑ Artificial Neural Networks (ANNs)
- ❑ ANNs components
- ❑ Forward Pass
- ❑ Training Process
- ❑ Cost Function
- ❑ Gradient Descent
- ❑ Backpropagation

What is regression supervised problem?

Example Predicting the score of a test based on the number of sleep and study hours on the night before.


Supervised: the examples have input and output.

Regression: test score is continuous value (if grade should be predicted then it is a classification problem).

Sleep	Study	Score
3	5	75
5	1	82
10	2	93
8	3	?


Biological motivation to ANNs

- ❑ ANNs architecture was inspired by how neurons in the brain are connected (see the figure below).
- ❑ The idea was to model human's brain in solving problems.
- ❑ A neuron is the basic computational unit in the brain.
- ❑ Synapses are the connections between neurons, it allows signals to flow all over the brain (signal-passing task).



Artificial Neural Networks


- ❑ ANNs architecture was inspired by how neurons in the brain are connected.
- ❑ The idea was to model human's brain in solving problems.
- ❑ A perceptron is the simplest form of ANNs, it contains number of inputs, number of outputs and only one neuron (node).
- ❑ Connections (synapses): $a_i = x_i * w_i$
- ❑ Neurons (nodes): $z = f(\sum x_i * w_i + b)$



Artificial Neural Networks Components

The architecture of this problem is as follows:

- ❑ 2 Inputs (Sleep and study hours).
- ❑ 1 Output (test score).
- ❑ 1 hidden layer with 3 nodes (neurons).
- ❑ Sigmoid activation function.



How to solve the problem


- ❑ Scaling $x_{norm} = x / \max(x)$.
- ❑ Define the architecture of the network (previous slide).
- ❑ Forward pass:
 - ❑ $a = x * w$
 - ❑ $b = f(a)$
 - ❑ $c = b * v$
 - ❑ $\hat{y} = f(c)$
- ❑ Train the network, this is done by minimizing the cost function: $J = \sum \frac{1}{2}(e_1^2 + e_2^2 + e_3^2)$ where $e = y - \hat{y}$.
- ❑ Numerical estimation to the final equation: $J = \sum \frac{1}{2}(y - f(f(x * w)v))^2$ Alternatively, finding the rate of change of J with respect to w, v : $\partial J / \partial w$ (This method called gradient descent).
- ❑ if $\partial J / \partial w$ is positive then the function is going up, otherwise, it is going down.
- ❑ The main reason we are using cost function as sum of squared errors to exploit the convex function nature.

Forward Pass

3	5
5	1
10	2

*

w_1^1	w_2^1	w_3^1
w_1^2	w_2^2	w_3^2



$3w_1^1 + 5w_1^2$	$3w_2^1 + 5w_2^2$	$3w_3^1 + 5w_3^2$
$5w_1^1 + 1w_1^2$	$5w_2^1 + 1w_2^2$	$5w_3^1 + 1w_3^2$
$10w_1^1 + 2w_1^2$	$10w_2^1 + 2w_2^2$	$10w_3^1 + 2w_3^2$

Training ANNs

- ❑ This step is followed by applying the sigmoid activation function. Then multiply the output with the weights of that connect hidden to output, lastly sigmoid activation function should be applied.
- ❑ $\partial J / \partial \mathbf{w} = \partial \sum \frac{1}{2}(\mathbf{y} - \hat{\mathbf{y}})^2 / \partial \mathbf{w} = \sum \partial \frac{1}{2}(\mathbf{y} - \hat{\mathbf{y}})^2 / \partial \mathbf{w}$
- ❑ The error at the last layer for single value: $\partial J / \partial \mathbf{w} = (\mathbf{y} - \hat{\mathbf{y}})$
- ❑ Back-propagation: chain rule, multiply the result by $\partial \hat{\mathbf{y}} / \partial \mathbf{w}$

Improve Results by applying Cost Function

- ❑ This step is followed by applying the sigmoid activation function. Then multiply the output with the weights of that connect hidden to output, lastly sigmoid activation function should be applied.
- ❑ $J = \frac{1}{2} \sum_{i=1}^N (y_i - \hat{y}_i)^2$
- ❑ We minimize the cost by changing the weights:

w_1^1	w_2^1	w_3^1
w_1^2	w_2^2	w_3^2

v_1^1
v_1^2
v_1^3

- ❑ So we can either apply "Semi-Brute Forcing" by checking range of possible values, alternatively, we can use Gradient Descent function.

Gradient Descent

- ❑ **Gradient Descent (GD)** is an algorithm that finds the weights (optimal weights) that make the cost function J the minimum.
- ❑ We can use either batch or stochastic GD (in this case we will use batch). In batch GD, we sum all the derivatives of J for all the observations:
 $\partial J / \partial w$.

Backpropagation

- ❑ **Backpropagation** is the technique used in training artificial neural network by updating the weights. Backpropagation make use of the GD algorithm.
- ❑ We have to compute 2 gradients: $\partial J / \partial \mathbf{w}$ and $\partial J / \partial \mathbf{v}$ the gradient with respect to the weight for hidden layer and the gradient with respect to the weight for output layer, respectively.

Backpropagation

- ❑ $\partial J / \partial \mathbf{v} = \mathbf{b}^T \delta_2$
 - ❑ $\delta_2 = -(y - \hat{y}) f'(c)$
 - ❑ $\partial J / \partial \mathbf{w} = \mathbf{x}^T \delta_1$
 - ❑ $\delta_1 = \delta_2 \mathbf{v}^T f'(a)$
-
- ❑ We should make our cost (J) as small as possible with an optimal combination of the weights (independently for both \mathbf{w} and \mathbf{v}).



Thank you :)